

Ansible Hands-On

DevOps Brasília

March 9, 2016

Deploy an application manually is usually time-consuming

- Most of the applications rely on different services to work correctly
- These services usually run on a distributed set of computing resources and communicate using various networking protocols
- Wire up these services by hand is time-consuming, error-prone, and it makes difficult to implement continuous delivery, for instance.
- Therefore, one way to deal with this problem is to use configuration management tools like Ansible, Chef, Puppet, Salt, among others.

What do we mean by configuration management?

- Writing the states for the servers, and then, using a tool to enforce that the servers are in the required state:
 - 1 the right packages are installed
 - 2 the configuration files contain the expected values and the correct permissions
 - 3 the right services are running
 - 4 ...

What can we expect from configuration management tools?

- they can help us on implementing continuous delivery, i.e., on implementing the blue-green deployment approach¹.
- on dealing with deployment orchestration. In other words, when there are multiple servers involved and the tasks must happen in a specific order. For instance, a database must be set up before bringing up the application servers.
- they can provision new servers. In the context of IaaS cloud, this means to spinning up new virtual machine instance.
- they help on guarantee the *idempotence* property.

And Ansible, what is it Good For?

- For describing the state of the servers through its DSL
- For doing deployment as well as configuration management
- For performing actions on multiple servers with a simple state model
- For control the order that the actions must happen in
- For talking to the public clouds API (e.g., AWS EC2, Google Compute Engine, Azure), as well as any cloud that supports the OpenStack API

How Ansible works




Figure 1: Using Ansible to perform actions on three remote servers

- ❶ it makes an SSH connection for each server
- ❷ it executes the first task on all the three nodes simultaneously:
 - ❶ generate a Python script that represents the task
 - ❷ copy the script to the servers
 - ❸ execute the script on the nodes
 - ❹ wait for the script to complete on all the hosts

What are good Ansible's characteristics?

- Easy-to-read syntax: its script (i.e., playbook) is built on top of the YAML format
- Nothing to install on the remote servers
- Push-based: it is the developer who controls when the changes happens to the servers
- Built-in modules: there are many modules to perform the tasks. Modules are idempotent
- A very thin layer of abstraction: we don't need to learn a new package manager
- It has a low learning curve




Figure 2: We will use Ansible to deploy WordPress with services running on two different nodes

Starting with Ansible

- 1 Provision the virtual machines and connect to the control node

```
vagrant up  
vagrant ssh
```

- 2 Create an inventory file to inform Ansible what are the remote servers, as well as how to connect to them.
- 3 Write the configuration scripts
- 4 Push the configuration scripts to the remote nodes

Ansible's inventory file

```
[defaults]
hostfile = hosts
remote_user = vagrant
private_key_file = ~/.ssh/id_rsa
host_key_checking = False
nocows = 1
```

```
[dbservers]
mysqlserver ansible_ssh_host=10.100.100.11 ansible_ssh_port=22
```

```
[webservers]
nginx ansible_ssh_host=10.100.100.12 ansible_ssh_port=22
```

- By default, Ansible looks for an inventory file (*ansible.cfg*) in:
 - 1 `ANSIBLE_CONFIG_ENVIRONMENT` variable
 - 2 `./ansible.cfg`
 - 3 `$HOME/.ansible.cfg`
 - 4 `/etc/ansible/ansible.cfg`

Ansible's hello world

```
ansible all -i hosts -m ping
```

or only

```
ansible all -m ping
```

where,

- **all** is the target hosts
- **hosts** is the host lists
- **ping** is the module (i.e., action) to execute

Default structure of an Ansible's project

/playbooks

— files

— static file₁

— static file₂

— ⋮

— static file_n

— templates

— template file₁.j2

— template file₂.j2

— ⋮

— template file_n

— ansible.cfg .2 hosts .2 playbook₁.yaml

— playbook₂.yaml

— ⋮

— playbook_n.yaml

Ansible's script: Playbook

Playbook

A playbook is the term that Ansible uses to describe a configuration management script. In practice, it is a list of plays.




Figure 3: Representing Ansible's playbook elements

Ansible's script: Playbook




Figure 4: Representing Ansible's playbook elements

What is a play?

A play associates an unordered set of hosts with an ordered list of tasks. A play must have:

- a set of hosts to configure
- a list of tasks to be executed in the hosts

Additionally, a play may also have:

- a **name**: a comment that describes what the play is about
- **become** and **become_method**: tell Ansible if the tasks must be executed as root
- **vars** a list of variables and values to be used in the play.

Ansible's script: Playbook




Figure 5: Representing Ansible's playbook elements

What is a task?

A task is the action to execute on the host. Every task must contain:

- a **key** with the name of a module
- a **value** with the arguments of the module.

```
tasks:
  - name: install nginx
    apt: name=nginx update_cache=yes cache_valid_time=3600
```

Playbooks support the usage of variables

- Variables can be used in tasks and in templates files.
- A variable's value can be: string, boolean, lists, and dictionaries.
- To reference a variable, we use the brace `{{ var_name }}` notation.
- Ansible uses Jinja2² template engine to evaluate variables in playbooks and in template files.

```
vars:
  cert_file: /etc/nginx/ssl/nginx.crt
tasks:
  - name: copy the TLS certificate
    copy: src=nginx.crt dest={{ cert_file }}
```

- Ansible also allows us to put the variables into one or more files, and reference them through the `vars_files` section.

```
vars_files:
  - nginx.yml
tasks:
  - name: copy the TLS certificate
    copy: src=nginx.crt dest={{ cert_file }}
```

Playbooks support the usage of variables

- Variables can be defined at runtime, using the **register** clause when executing a task.
- In this case, the type of the variable is always a dictionary, and its keys depend of the modules.
- the debug task can be used to know the value of a variable.

```
tasks:  
  - name: register the output of whoami command  
    command: whoami  
    register: login_user  
  - debug: msg="Logged as user: {{ login_user.stdout }}"
```

Built-in variables

Name	Description
hostvars	a dictionary whose keys are Ansible hostnames and values are dictionaries that map variables names to values.
inventory_hostname	name of the current host as defined in the inventory file.
group_names	a list of all groups that the current node is member of
groups	a dictionary whose keys are Ansible group names and values are list of hostnames that are member of the group.
play_hosts	a list of the hostnames of the current play.
ansible_version	a dictionary with the Ansible's version.

- Ansible can also provide information about the node, such as IP addresses, memory size, disk, operating system, etc.

```
---
- name: collect the name of the user and the facts about the
  node
  hosts: webservers
  gather_facts: yes
  tasks:
    - name: register the output of whoami command
      command: whoami
      register: login_user
    - debug: msg="Logged as user {{ login_user.stdout }} and my
      IP address is {{ hostvars[groups['webservers'][0]]['
      ansible_eth1']['ipv4']['address'] }}"
```

Using Handler to notify a new state

- A handler is similar to a tasks, but it only runs if it has be notified by a task.
- A task only fires a notification only if the node's state has changed.
- Handlers only run after all of the tasks have finished, and they only run once, even if they are notified multiple times.
- Handlers always run in the order that they appear in the play, and not in the notification order.

```
tasks:
  - name: copy the TLS key
    copy: src=files/nginx.key dest={{ key_file }} owner=root
        mode=0600
    notify: restart nginx
handlers:
  - name: restart nginx
    service: name=nginx state=restarted
```
